

INSTALL (all-in-one) English

Preparation [↗](#)

Hardware [↗](#)

This is the minimum hardware configuration for running all-in-one as follows. The hardware names in parentheses are those used in the test environment.

When you use different hardware below, you will need to read the steps again and change the settings.

- 8 CPU Cores
- 32GB RAM
- 100GB hard disk space
- 1x Ethernet NIC
- 2x SR-IOV PF(Dual Port Mellanox ConnectX-5)
- 1x NVIDIA GPU(A100)

Software [↗](#)

Create an environment with the following installed. Set PATH for kind, go, and kubectl.

- Ubuntu
 - 22.04.1
- Kernel version
 - 5.15.0-122-generic or higher
- kind
 - 0.23.0
- Docker
 - 26.1.3, build b72abbb
- go
 - 1.23.0
- Mellanox OFED Driver
 - MLNX_OFED_LINUX-24.04-0.7.0.0-ubuntu22.04-x86_64.tgz
- kubectl
 - 1.31.1-1.1

Network design [↗](#)

Design the IP addresses for each segment in advance. The harbor server is described as an example of a local container registry.

	External IP of the Pod	172.18.9.0-172.18.9.10	./manifest/metallb/ipaddrespool.yaml	<ul style="list-style-type: none"> IP address range that MetalLB assigns to the Service Change according to the Docker Node IP segment
Kubernetes	PodSubnet	10.100.0.0/16	./kind/kind-values/values.yaml	<ul style="list-style-type: none"> Specify CIDR to be used for CNI communication Set an unused one in the environment

Additional host settings [↗](#)

Disable swap [↗](#)

Delete swap space as needed.

```
1 $ sudo swapoff -a
2 $ sudo swapon --show
```

Configure Hugepages [↗](#)

1. Edit /etc/default/grub

```
1 $ sudo vi /etc/default/grub
2 ...
3 GRUB_CMDLINE_LINUX_DEFAULT="default_hugepagesz=1G hugepagesz=1G hugepages=32"
```

2. Apply Grub settings and reboot

```
1 $ sudo update-grub
2 $ sudo reboot
```

3. Check Hugepage settings

```
1 $ cat /proc/meminfo
2 :
3 HugePages_Total: 32
4 HugePages_Free: 32
5 HugePages_Rsvd: 0
6 HugePages_Surp: 0
7 Hugepagesize: 1048576 kB
```

Enable SR-IOV feature [↗](#)

Here is an example of Mellanox ConnectX-5

1. Get the driver below

- [https://developer.nvidia.com/networking/mlnx-ofed-eula?](https://developer.nvidia.com/networking/mlnx-ofed-eula?mtag=linux_sw_drivers&mrequest=downloads&mtype=ofed&mver=MLNX_OFED-24.04-0.7.0.0&mname=MLNX_OFED_LINUX-)
mtag=linux_sw_drivers&mrequest=downloads&mtype=ofed&mver=MLNX_OFED-24.04-0.7.0.0&mname=MLNX_OFED_LINUX-

24.04-0.7.0.0-ubuntu22.04-x86_64.tgz

- Files used

```
1 $ md5sum MLNX_OFED_LINUX-24.04-0.7.0.0-ubuntu22.04-x86_64.tgz
2 8a15626083c14dcd64e33f200ebd142a MLNX_OFED_LINUX-24.04-0.7.0.0-ubuntu22.04-x86_64.tgz
```

2. Install driver

- Extract tar

```
1 $ sudo su
2 # tar xzvf MLNX_OFED_LINUX-24.04-0.7.0.0-ubuntu22.04-x86_64.tgz
3 ./MLNX_OFED_LINUX-24.04-0.7.0.0-ubuntu22.04-x86_64/
4 ./MLNX_OFED_LINUX-24.04-0.7.0.0-ubuntu22.04-x86_64/DEBS/
5 ...
6 ./MLNX_OFED_LINUX-24.04-0.7.0.0-ubuntu22.04-x86_64/.arch
7 ./MLNX_OFED_LINUX-24.04-0.7.0.0-ubuntu22.04-x86_64/distro
```

- Run installer

```
1 cd MLNX_OFED_LINUX-24.04-0.7.0.0-ubuntu22.04-x86_64
2
3 # ./mlnxofedinstall
4 Logs dir: /tmp/MLNX_OFED_LINUX.22943.logs
5 General log file: /tmp/MLNX_OFED_LINUX.22943.logs/general.log
6
7 Below is the list of MLNX_OFED_LINUX packages that you have chosen
8 (some may have been added by the installer due to package dependencies):
9
10 ofed-scripts
11 ...
12 ibarr
13
14 This program will install the MLNX_OFED_LINUX package on your machine.
15 Note that all other Mellanox, OEM, OFED, RDMA or Distribution IB packages will be removed.
16 Those packages are removed due to conflicts with MLNX_OFED_LINUX, do not reinstall them.
17
18 Do you want to continue?[y/N]:y
19
20 Checking SW Requirements...
21 One or more required packages for installing MLNX_OFED_LINUX are missing.
22 ...
23 Installation passed successfully
24 To load the new driver, run:
25 /etc/init.d/openibd restart
```

3. Load driver

```
1 # /etc/init.d/openibd restart
2 Unloading HCA driver: [ OK ]
3 Loading HCA driver and Access Layer: [ OK ]
```

4. Enable SR-IOV

- Survey of target PF

```
1 # ip link show ens9f0np0
2 8: ens9f0np0: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN mode DEFAULT group default qlen 1000
3    link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff
4    altname enp23s0f0np0
```

```

5 # ip link show ens9f1np1
6 9: ens9f1np1: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN mode DEFAULT group default qlen 1000
7     link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff
8     altname enp23s0f1np1
9
10 lspci -Dnn | grep Mellanox
11 0000:17:00.0 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex] [15b3:1019]
12 0000:17:00.1 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex] [15b3:1019]

```

- Run the SR-IOV enable command

```

1 # mstconfig -d 0000:17:00.0 set SRIOV_EN=1 NUM_OF_VFS=8
2 # mstconfig -d 0000:17:00.1 set SRIOV_EN=1 NUM_OF_VFS=8

```

- Reboot host

```

1 # reboot

```

- Create VFs

```

1 # echo 8 > /sys/class/net/ens9f0np0/device/sriov_numvfs
2 # echo 8 > /sys/class/net/ens9f1np1/device/sriov_numvfs

```

- Check VFs

```

1 # ip addr
2 ...
3 5: ens9f0np0: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
4     link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff
5     altname enp23s0f0np0
6 6: ens9f1np1: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
7     link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff
8     altname enp23s0f1np1
9 7: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN group default
10     link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff
11     inet 172.17.0.1/16 brd 172.17.255.255 scope global docker0
12         valid_lft forever preferred_lft forever
13 8: ens9f0v0: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
14     link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff permaddr xx:xx:xx:xx:xx:xx
15     altname enp23s0f0v0
16 9: ens9f0v1: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
17     link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff permaddr xx:xx:xx:xx:xx:xx:
18     altname enp23s0f0v1
19 10: ens9f0v2: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
20     link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff permaddr xx:xx:xx:xx:xx:xx:
21     altname enp23s0f0v2
22 11: ens9f0v3: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
23     link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff permaddr xx:xx:xx:xx:xx:xx:
24     altname enp23s0f0v3
25 12: ens9f0v4: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
26     link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff permaddr xx:xx:xx:xx:xx:xx:
27     altname enp23s0f0v4
28 13: ens9f0v5: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
29     link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff permaddr xx:xx:xx:xx:xx:xx:
30     altname enp23s0f0v5
31 14: ens9f0v6: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
32     link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff permaddr xx:xx:xx:xx:xx:xx:
33     altname enp23s0f0v6
34 15: ens9f0v7: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
35     link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff permaddr xx:xx:xx:xx:xx:xx:

```

```

36  altname enp23s0f0v7
37  16: ens9f1v0: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
38      link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff permaddr xx:xx:xx:xx:xx:
39      altname enp23s0f1v0
40  17: ens9f1v1: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
41      link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff permaddr xx:xx:xx:xx:xx:
42      altname enp23s0f1v1
43  18: ens9f1v2: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
44      link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff permaddr xx:xx:xx:xx:xx:
45      altname enp23s0f1v2
46  19: ens9f1v3: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
47      link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff permaddr xx:xx:xx:xx:xx:
48      altname enp23s0f1v3
49  20: ens9f1v4: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
50      link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff permaddr xx:xx:xx:xx:xx:
51      altname enp23s0f1v4
52  21: ens9f1v5: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
53      link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff permaddr xx:xx:xx:xx:xx:
54      altname enp23s0f1v5
55  22: ens9f1v6: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
56      link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff permaddr xx:xx:xx:xx:xx:
57      altname enp23s0f1v6
58  23: ens9f1v7: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
59      link/ether xx:xx:xx:xx:xx:xx brd ff:ff:ff:ff:ff:ff permaddr xx:xx:xx:xx:xx:
60      altname enp23s0f1v7
61
62  # lspci -Dnn | grep Mellanox | grep Virtual
63  0000:17:00.2 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex Virtual Function] [15b3:101a]
64  0000:17:00.3 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex Virtual Function] [15b3:101a]
65  0000:17:00.4 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex Virtual Function] [15b3:101a]
66  0000:17:00.5 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex Virtual Function] [15b3:101a]
67  0000:17:00.6 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex Virtual Function] [15b3:101a]
68  0000:17:00.7 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex Virtual Function] [15b3:101a]
69  0000:17:01.0 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex Virtual Function] [15b3:101a]
70  0000:17:01.1 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex Virtual Function] [15b3:101a]
71  0000:17:01.2 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex Virtual Function] [15b3:101a]
72  0000:17:01.3 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex Virtual Function] [15b3:101a]
73  0000:17:01.4 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex Virtual Function] [15b3:101a]
74  0000:17:01.5 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex Virtual Function] [15b3:101a]
75  0000:17:01.6 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex Virtual Function] [15b3:101a]
76  0000:17:01.7 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex Virtual Function] [15b3:101a]
77  0000:17:02.0 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex Virtual Function] [15b3:101a]
78  0000:17:02.1 Ethernet controller [0200]: Mellanox Technologies MT28800 Family [ConnectX-5 Ex Virtual Function] [15b3:101a]

```

Placement of necessary files

Place the following files in the repository dir

Location	File name	Note
./src/	controller.tar.gz	<ul style="list-style-type: none"> Clone the following repository https://github.com/openkasugai/controller Archive the repository into a tar.gz archive

		<ul style="list-style-type: none"> • Change directory name to <code>controller</code> • Set recursive option when cloning repository
<code>./driver/nvidia/</code>	NVIDIA-Linux-x86_64-550.90.12.run	<ul style="list-style-type: none"> • Get it from below https://www.nvidia.com/Download/index.aspx
<code>./work/</code>	*.mp4	<ul style="list-style-type: none"> • Place video files used to run video inference scenarios <ul style="list-style-type: none"> ◦ Video files (*.mp4) should be located in <code>./work/DATA/video</code>. • This directory will bind to <code>/root/work</code> in each Docker container

Makefile creation procedure [↗](#)

Modify the parameters in the Makefile to suit your build environment.

Parameter Description [↗](#)

Parameter Name	Description	Default
DCI_REGISTRY_ADDR	OpenKasugai Controller Registry URL	<code>ghcr.io/openkasugai/controller</code>
DCI_REGISTRY_ADDR_ALT	Registry URL for the harbor server	<code>192.168.10.100/images</code>
DCI_REGISTRY_CERT_URL	Download URL of the registry certificate file of the harbor server	<code>http://192.168.10.100/harbor/ca.crt</code>
DCI_K8S_SOFTWARE_TARFILE	The file name of the tar-compressed OpenKasugai controller code	<code>controller.tar.gz</code>
NVIDIA_DRIVER_FILE	NVIDIA driver install script file name	<code>NVIDIA-Linux-x86_64-550.90.12.run</code>
K8S_VERSION	The version of the k8s cluster to be built with kind	<code>v1.31.0</code>
GO_VERSION	Go version to install on kind node	<code>1.23.0</code>
SRIOV_CNI_VERSION	The version of SRIOV to use	<code>v2.8.1</code>
CRIO_OS	The OS version to use for the kind node	<code>xUbuntu_22.04</code>
CRIO_REPO_VERSION	The version of cri-o to install on the k8s cluster (repository)	<code>v1.31</code>
CRIO_PKG_VERSION	The version of cri-o to install on the k8s cluster (package)	<code>1.31.0-1.1</code>
METALLB_VERSION	The version of MetalLB to use	<code>v0.14.8</code>

NVIDIA_K8S_IPAM_VERSION	The version of nvidia-k8s-ipam to use	v0.3.5
IMG_TAG	Image tag name when retrieving an image from DockerHub	22.04.5
SRIOV_IF_REGEX_1	Regular expression to get the IF names of both the first PF and VF to be assigned to kind-worker	ens9f0.*
SRIOV_IF_REGEX_2	Regular expression to get the IF names of both the second PF and VF to be assigned to kind-worker	ens9f1.*
SRIOV_PF_IF_1	IF name of the first PF to be assigned to kind-worker	ens9f0np0
SRIOV_PF_IF_2	IF name of the second PF to be assigned to kind-worker	ens9f1np1
SRIOV_PF_IPADDR_1	IP address (CIDR) to be assigned to the first PF	192.168.20.2/24
SRIOV_PF_IPADDR_2	IP address (CIDR) to be assigned to the second PF	192.168.20.3/24
SRIOV_NUM_VFS	Number of VFs to be created	8
NV_IPAM_SUBNET	Subnet of IP address automatically assigned to VF	192.168.20.0/24
NV_IPAM_PER_NODE_BLOCK_SIZE	Number of IP addresses per node for dividing IP addresses automatically assigned to VFs	30
NV_IPAM_GATEWAY	Gateway IP address automatically assigned to the VF	192.168.20.1
NV_IPAM_EXCLUDE_START_1	Start of IP address to be excluded (first half) from IP addresses automatically assigned to VFs	192.168.20.0
NV_IPAM_EXCLUDE_END_1	End of IP address to be excluded (first half) from IP addresses automatically assigned to VFs	192.168.20.9
NV_IPAM_EXCLUDE_START_2	Start of IP address to be excluded (second half) from IP addresses automatically assigned to VFs	192.168.20.51
NV_IPAM_EXCLUDE_END_2	End of IP address to be excluded (second half) from IP addresses automatically assigned to VF	192.168.20.255

Explanation of main commands

This section describes the main commands accepted by Makefile.

Refer to [README.md](#) for instructions on how to build a cluster.

Commands	Description
make images	Run the image build.
make create-cluster-with-all	Starting the OpenKasugai controller cluster.
make dataflow-<PatternName>	Deploying a Scenario.

Image Creation Procedure

Run below command and create the container images used by kind-control-plane, kind-worker, and kind-worker2.

```
1 make images
```

The images created above can be used for the following purposes.

Image	Parent image	Description	kind-control-plane	kind-worker ※with accelerator	kind-worker2 ※without accelerator
kind-ubuntu	ubuntu official image	Base image for kind node, determines the OS version.	—	—	—
kind-ubuntu-node	kind-ubuntu	The image with k8s etc installed on base.	—	—	—
kind-ubuntu-node-crio	kind-ubuntu-node	The above image, installed crio as container runtime.	—	—	—
dci-kind-node-non-acc	kind-ubuntu-node-crio	Image for nodes without GPUs	○	—	○
dci-kind-node-with-acc	dci-kind-node-non-acc	Image for nodes with GPUs NVIDIA driver installation is done internally.	—	○	—

Cluster construction procedure

Execute one of the following. Select the latter if you want to learn the role of OpenKasugai-controller resources or run a scenario you created yourself.

- Bulk build procedure
- Procedure for manually registering resources

Bulk build procedure

By executing the following, all the settings required to run the sample scenario will be configured.

```
1 make create-cluster-with-all
```

Procedure for manually registering resources

If you want to manually register the resources required for scenario execution later, do the following:

```
1 make create-cluster-without-senario
```

Cluster usage procedure [↗](#)

Running kubectl from the host [↗](#)

The `.kube/config` of the cluster created by kind is placed in the home directory of the root user. To use kubectl, run it with root privileges as follows, or copy `.kube/config` to the home directory of any user.

```
1 sudo kubectl get nodes
```

Scenario execution [↗](#)

- See [RUN_SCENARIOS](#)

Cluster deletion procedure [↗](#)

By executing the following, the OpenKasugai-controller execution environment will be deleted.

```
1 make delete-cluster
2 make clean
```